
Supplementary material: MAHA

Anonymous Author(s)

Affiliation

Address

email

1 A Explanation on Set Transformer

2 Set Transformer [14] is proven to be a flexible function approximator that considers a high-order
3 interaction between set elements. It can be decomposed into the following 4 attention modules:

$$\text{MAB}(A, B) = \text{LN}(H + \text{rFF}(H)) \in \mathbb{R}^{n \times d}$$

$$\text{SAB}(A) = \text{MAB}(A, A) \in \mathbb{R}^{n \times d}$$

$$\text{ISAB}_m(A) = \text{MAB}(A, \text{MAB}(I, A)) \in \mathbb{R}^{n \times d}$$

$$\text{PMA}_k(A) = \text{MAB}(S, \text{rFF}(A)) \in \mathbb{R}^{k \times d}$$

4 where $H = \text{LN}(A + \text{Multihead}(A, B, B))$ is a basic building block for every module. Here,
5 $A, B \in \mathbb{R}^{n \times d}$ are random sets, and $I \in \mathbb{R}^{m \times d}, S \in \mathbb{R}^{k \times d}$ are additional learnable parameters. Note
6 that the randomly initialized inducing points I in ISAB have a lower cardinality than A .

7 A multi-head attention block (MAB) and set attention block (SAB) are the two main key components,
8 which reinforce the multi-head-attention and self-attention with a layer normalization and a skip
9 connection. The induced set attention block (ISAB), with m inducing points, is further devised as a
10 substitute for the SAB in terms of computational efficiency and generalization. The output size is
11 fixed to k by another complex pooling module: pooling by multi-head attention (PMA). Throughout
12 the experiments, 2 ISAB, 1 PMA, 2 SAB are composited in order. m is set to 32 in regression, 256 in
13 classification, and k is set to 1 imitating the mean-pooling operation over *shot* in NP.

14 B Distilling an obtainable knowledge from T to C

15 Meta-learning has shown to be vulnerable to overfitting due to task ambiguity [19, 33, 16, 36, 32].
16 According to [20], there are too many local optima in meta-learning that can lead to bad test results
17 when it comes to a limited number of both *way* and *shot*. Here, we mainly deal with the few-shot
18 nature of meta-learning such that additional regularization terms are devised to distill an obtainable
19 knowledge from T to C .

20 Since the output distribution accounts for a significant portion of the variability of the neural processes
21 [12], we minimize KL divergence between the following output distributions:

$$o(T_y|C) := \mathbb{E}_{q(r|C)} [p(T_y|T_x, r, z)], \quad o(T_y|T) := \mathbb{E}_{q(r|T)} [p(T_y|T_x, r, z)]$$

22 Notice that the deterministic representations are conditioned on the different sets, one from the context
23 set C , another from the target set T , and we assume that the stochastic representations are given in
24 advance. For a general-purpose, we derive an upper bound as follows since the KL divergence can be
25 computed in a closed-form only in a limited family of probability distributions:

$$\begin{aligned} KL(o(T_y|C) \| o(T_y|T)) &= - \int o(T_y|C) \log o(T_y|T) dT_y - \mathcal{H}(o(T_y|C)) \\ &\approx - \log o(\hat{T}_y|T) - \mathcal{H}(o(T_y|C)) \quad \text{s.t.} \quad \hat{T}_y \sim o(T_y|C) \\ &\leq - \mathbb{E}_{q(r|T)} [\log p(\hat{T}_y|T_x, r, z)] - \mathcal{H}(o(T_y|C)) \end{aligned}$$

where $\mathcal{H}(\cdot)$ indicates entropy. The approximation is conducted using a Monte Carlo sample, and the inequality is from Jensen’s inequality on the concave $\log(\cdot)$ function. The first term in the last line is conceptually similar to cross entropy, which leads the model prediction to refer to the pseudo-label which we detach from the computational graph to avoid cycle following [15]. The second term helps the model to avoid overconfidence and degeneracy as discussed in [25, 8, 10]. As a result, the loss function can be rewritten as follow by augmenting the regularization term :

$$\begin{aligned} \mathcal{L}_{KD} = & -\mathbb{E}_{q(r|C)q(z|T)} [\log p(T_y|T_x, r, z)] + \beta_3 \cdot KL(q(z|T)||q(z|C)) \\ & - \beta_3 \cdot \left(\mathbb{E}_{q(r|T)q(z|T)} \left[\log p(\hat{T}_y|T_x, r, z) \right] + \mathbb{E}_{q(z|T)} [\mathcal{H}(o(T_y|C))] \right) \end{aligned} \quad (1)$$

Note that it can bridge to studies on knowledge distillation, specifically a self-distillation [34, 9, 30, 13], where a network is trained not only with the true output T_y , but also with the soft output \hat{T}_y that is estimated by the network itself.

C Implementation details

C.1 Dataset

Gaussian Process A batch of size 16, context set of variable size ranged from 5 to 10, and target set of size 30 are considered. For the squared exponential kernel $k(x, x') = \sigma^2 \exp(-0.5(x - x')^2/l^2)$, the hyperparameters are chosen to be $l = 0.5$ and $\sigma = 1$. Inputs are uniformly sampled from $[-2.0, 2.0]$ and outputs are computed based on the Cholesky decomposition of the kernel with the noise parameter $\sigma_n = 0.02$ [21].

Sine&Polynomial A batch of size 25, context set of size 5 or 10, and target set of size 15 or 20 are considered. Input domain is fixed to $[-5.0, 5.0]$, and a task is defined among the four functions whose coefficients are uniformly sampled from the intervals summarized in Table 1.

Table 1: Coefficient settings

	Sine	Line	Quad	Cubic
A	[0.1, 5.0]	[-3.0, 3.0]	[-0.2, 0.2]	[-0.1, 0.1]
B	[0.8, 1.2]	[-3.0, 3.0]	[-2.0, 2.0]	[-0.2, 0.2]
C	[0.0, 2π]	-	[-3.0, 3.0]	[-2.0, 2.0]
D	-	-	-	[-3.0, 3.0]

Mini-ImageNet, Tiered-ImageNet A batch of size 12 is considered where each batch instance is generated by sampling five random classes from the meta set with randomly assigned labels from $\{0, 1, 2, 3, 4\}$. Then, for each of the chosen classes, 1 or 5 images are selected as the context set, and 15 other images are additionally selected to construct the target set.

Multi-dataset Task generation process and size of the context/target set are equal to the setting in mini-ImageNet and tiered-ImageNet. However, unlike mini-ImageNet or tiered-ImageNet, images are not pre-processed in advance by the deep residual network. Instead, all images in the meta-train set, meta-valid set, and meta-test set are resized to $84 \times 84 \times 3$, and Conv-blocks are utilized to extract the feature from the images. Due to extensive memory usage during the feature extraction, a small batch of size 4 is considered where each batch instance is generated among the four fine-grained image classification datasets.

Table 2: Summary of classification datasets

Dataset	mini-ImageNet	tiered-ImageNet	Bird	Texture	Aircraft	Fungi
Source	[23]	[23]	[28]	[7]	[18]	[1]
Split setting	[26]	[22]	[31]	[31]	[31]	[31]
Finess	Coarse	Coarse	Fine	Fine	Fine	Fine

56 C.2 Architecture design

57 We show the detailed architectures used for the feature extractor in Table 3. Here, Conv(d, k, s, n,
58 p) is a convolutional block with d output channels, k kernel size, s stride size, n normalization, p
59 pooling method. LRN and BN indicate a local response normalization and a batch normalization,
60 respectively, and MAX is a max-pooling with a kernel size 3 and stride 2. By default, two linear
61 layers are commonly exploited, which come after the convolutional layers if the model input is a 3D
62 image. For the convolutional layers, we follow the exact setting of [31] depending on whether the
63 model is for task clustering or prediction. For the dropout rate p , please refer to Section C.3.

Table 3: Feature extractor $g(\cdot)$ architecture

Gaussian Process Sine&Polynomial	mini-ImageNet tiered-ImageNet	multi-dataset	
		2 Conv	4 Conv
Linear(1, 128)	Dropout(p)	Conv(32, 5, 1, LRN, MAX)	Conv(32, 3, 1, BN, MAX)
ReLU	Linear(640, 128)	Conv(32, 5, 1, LRN, MAX)	Conv(32, 3, 1, BN, MAX)
Linear(128, 128)	ReLU	Linear($32 \times 21 \times 21$, 384)	Conv(32, 3, 1, BN, MAX)
	Linear(128, 128)	ReLU	Conv(32, 3, 1, BN, MAX)
		Linear(384, 128)	Dropout(p)
			Linear($32 \times 5 \times 5$, 128)
			ReLU
			Linear(128, 128)

64 In Table 4 and 5, the encoder-decoder pipeline of MAHA is summarized. Note that the encoders for r
65 and z are almost the same except for the output size, which is doubled in z due to reparameterization.
66 Also, note that the size of the inputs is different between regression and classification. This is because
67 $g(X)$ and Y are concatenated in regression while $shots$ of $g(X)$ is first divided along way by Y and
68 then separately feed-forwarded in classification. Lastly, in regression, the encoder outputs, r and
69 (reparameterized) z , are reshaped from from $[batch, 1, 256]$ into $[batch, 2, 128]$ before feed-forwarded
70 into the decoder. By default, all networks use the Adam optimizer with a constant learning rate and
71 an l2 regularization of weight $1e-4$. Please refer to the attached code for actual implementation.

Table 4: Regression architecture

Encoder		Decoder
Enc _{r} (\cdot)	Enc _{z} (\cdot)	rFF(\cdot)
ISAB ₃₂ (128 + 1, 128)	ISAB ₃₂ (128 + 1, 128)	Linear(128, 128)
ISAB ₃₂ (128, 128)	ISAB ₃₂ (128, 128)	ReLU
PMA ₁ (128, 128)	PMA ₁ (128, 128)	Linear(128, 128)
SAB(128, 128)	SAB(128, 128)	
SAB(128, 256)	SAB(128, 512)	

Table 5: Classification architecture

Encoder		Decoder
Enc _{r} (\cdot)	Enc _{z} (\cdot)	rFF(\cdot)
ISAB ₂₅₆ (128, 128)	ISAB ₂₅₆ (128, 128)	Linear(128, 128)
ISAB ₂₅₆ (128, 128)	ISAB ₂₅₆ (128, 128)	ReLU
PMA ₁ (128, 128)	PMA ₁ (128, 128)	Linear(128, 128)
SAB(128, 128)	SAB(128, 128)	
SAB(128, 128)	SAB(128, 256)	

72 C.3 Hyperparameter

73 Hyperparameters are optimized with the validation loss of the model trained on the train meta-set.
 74 Among the many hyperparameter optimization processes [3, 4, 17, 29], we use the random search
 75 whose outcomes are summarized in Table 6 and 9. For the heterogeneous datasets, an agglomerative
 76 clustering is applied to the t-SNE embeddings of $\mu_{\bar{z}}$ from the stochastic path where we use the default
 77 setting of Scipy [27], an open-source scientific tools for Python. In Table 7 to 8 and Table 10 to 11,
 78 the clustering results for randomly generated 2500 (in regression) or 4000 (in classification) data
 79 points are presented by cross-tabulation between the cluster index and the true dataset label. For the
 80 reason why only two clusters are considered in Sine&Polynomial, please refer to Section D.

Table 6: Hyperparameters for regression

Gaussian Process		Sine&Polynomial			
		5-shot		10-shot	
epoch (pre)	-	1e+6		1e+6	
lr (pre)	-	1e-4		1e-4	
β_1	-	1		1	
		Cluster index			
		1st	2nd	1st	2nd
epoch	1e+6	1e+6	1e+6	1e+6	1e+6
lr	1e-4	1e-4	1e-4	1e-4	1e-4
β_2	1	1	1	1	1
β_3	-	0.120	0.941	0.156	0.224

Table 7: 5-shot Sine&Polynomial

Cluster index	Sine	Line	Quad	Cubic
1st	584	9	10	7
2nd	11	589	614	676

Table 8: 10-shot Sine&Polynomial

Cluster index	Sine	Line	Quad	Cubic
1st	587	0	0	0
2nd	1	598	658	656

Table 9: Hyperparameters for classification

	mini-ImgeNet		tiered-ImgeNet		multi-dataset							
	1-shot	5-shot	1-shot	5-shot	1-shot				5-shot			
epoch (pre)	-	-	-	-	1.4e+4				1.4e+4			
lr (pre)	-	-	-	-	4.1e-5				4.3e-5			
β_1	-	-	-	-	0.017				0.016			
					Cluster index							
					1st	2nd	3rd	4th	1st	2nd	3rd	4th
epoch	5e+4	6e+4	6e+4	7e+4	2e+4	1.5e+4	5e+3	2e+4	5e+3	1e+4	1e+4	1e+4
lr	8.7e-5	9.2e-5	8.8e-5	9.9e-5	4.4e-5	3.0e-5	7.8e-5	2.5e-5	8.1e-5	7.1e-5	2.9e-5	7.9e-5
p	0.3	0.3	0.47	0.41	0.4	0.4	0.4	0.4	0.4	0.4	0.4	0.4
β_2	0.098	0.004	0.091	0.001	0.063	0.099	0.064	0.020	0.032	0.050	0.032	0.075
β_3	-	-	-	-	0.085	0.068	0.063	0.051	0.018	0.085	0.028	0.053

Table 10: 1-shot multi-dataset

Cluster index	Bird	Texture	Aircraft	Fungi
1st	9	0	950	0
2nd	930	17	45	27
3rd	11	12	0	921
4th	6	1023	0	49

Table 11: 5-shot multi-dataset

Cluster index	Bird	Texture	Aircraft	Fungi
1st	1030	0	0	0
2nd	0	0	0	979
3rd	0	1006	0	3
4th	0	0	982	0

81 C.4 Computing Resource

82 We run experiments on NVIDIA GeForce RTX 2080 Ti, which takes two days for regression and three
 83 days for classification. Compared to the basic neural processes, MAHA is about two times slower
 84 in terms of the convergence speed, mainly due to the flexible encoder, Set Transformer. However,
 85 the compatibility and necessity are empirically verified by the outstanding performance of both
 86 homogeneous and heterogeneous datasets. Many follow-up studies are emerging these days to speed
 87 up the training of Transformers when applying the attention-based modules [11, 6], which would
 88 make our work more valid.

89 D Additional experimental results

90 D.1 Clustering result

91 In Table 1, note that Quad is perfectly covered by Cubic, and Quad and Cubic mostly cover Line.
 92 Hence, rather than four separate clusters, only two are shown in Figure 1, each of which implies either
 93 sine or polynomial. On the other hand, in Figure 2, four distinct fine-grained image classification
 94 datasets are clearly discriminated by separate clusters.




Figure 1: t-SNE visualization of the task representation from Sine&Polynomial




Figure 2: t-SNE visualization of the task representation from multi-dataset

95 D.2 POOL resolves KL collapse

96 Although a small subset C of T is expected to reproduce
 97 the stochastic representation through the KL divergence
 98 in the loss function, the representation inferred by T is
 99 restricted to be underutilized, which is the KL collapse
 100 [5, 2, 24, 35]. By dimension-wise pooling operations, we
 101 intended to prevent z from being redundant by allowing
 102 the information flow to go through the stochastic path
 103 whenever heterogeneous tasks occur in *batch*. In Figure 3,
 104 the training curve of the KL divergence is visualized for
 105 many different encoder-decoder pipelines: NP, NP+LD,
 106 NP+FE, FELD. It can be observed that the KL collapse no
 107 more occurs when accompanied by the pooling operations,
 108 which implies that the posterior does not simply converge
 109 to the prior.




Figure 3: Moving average of KL divergence along epochs

110 D.3 POOL needs AE for interpretability

111 Although the pooling operations are applied to disentangle the task representation from the stochastic
 112 path, we observe in the paper that the auto-encoding structure is additionally required to achieve
 113 interpretability and high purity values. It is mainly due to the restricted flexibility of \bar{r} , which
 114 encourages \bar{z} to imply not only the heterogeneity but also the local features that are initially in
 115 charge of the deterministic path. Here, the auto-encoding structure allows r to be inferred by the
 116 (large) target set T , not the (small) context set C , which is advantageous to obtain a more flexible
 117 set representation. Therefore, the restricted flexibility of \bar{r} can be resolved so that \bar{z} can provide
 118 well-clustered and interpretable task representation.

References

- [1] 2018 fgcvx fungi classification challenge. 2018.
- [2] Alexander A Alemi, Ian Fischer, Joshua V Dillon, and Kevin Murphy. Deep variational information bottleneck. *arXiv preprint arXiv:1612.00410*, 2016.
- [3] James Bergstra and Yoshua Bengio. Random search for hyper-parameter optimization. *Journal of machine learning research*, 13(2), 2012.
- [4] James Bergstra, Daniel Yamins, and David Cox. Making a science of model search: Hyperparameter optimization in hundreds of dimensions for vision architectures. In *International conference on machine learning*, pages 115–123. PMLR, 2013.
- [5] Samuel R Bowman, Luke Vilnis, Oriol Vinyals, Andrew M Dai, Rafal Jozefowicz, and Samy Bengio. Generating sentences from a continuous space. *arXiv preprint arXiv:1511.06349*, 2015.
- [6] Krzysztof Choromanski, Valerii Likhoshesterov, David Dohan, Xingyou Song, Andreea Gane, Tamas Sarlos, Peter Hawkins, Jared Davis, Afroz Mohiuddin, Lukasz Kaiser, et al. Rethinking attention with performers. *arXiv preprint arXiv:2009.14794*, 2020.
- [7] Mircea Cimpoi, Subhansu Maji, Iasonas Kokkinos, Sammy Mohamed, and Andrea Vedaldi. Vedaldi a.: Describing textures in the wild. In *In Computer Vision and Pattern Recognition (CVPR)*. Citeseer, 2014.
- [8] Tuomas Haarnoja, Aurick Zhou, Pieter Abbeel, and Sergey Levine. Soft actor-critic: Off-policy maximum entropy deep reinforcement learning with a stochastic actor. In *International Conference on Machine Learning*, pages 1861–1870. PMLR, 2018.
- [9] Yuenan Hou, Zheng Ma, Chunxiao Liu, and Chen Change Loy. Learning lightweight lane detection cnns by self attention distillation. In *Proceedings of the IEEE/CVF International Conference on Computer Vision*, pages 1013–1021, 2019.
- [10] Xu Ji, João F Henriques, and Andrea Vedaldi. Invariant information clustering for unsupervised image classification and segmentation. In *Proceedings of the IEEE International Conference on Computer Vision*, pages 9865–9874, 2019.
- [11] Nikita Kitaev, Łukasz Kaiser, and Anselm Levskaya. Reformer: The efficient transformer. *arXiv preprint arXiv:2001.04451*, 2020.
- [12] Tuan Anh Le, Hyunjik Kim, Marta Garnelo, Dan Rosenbaum, Jonathan Schwarz, and Yee Whye Teh. Empirical evaluation of neural process objectives. In *NeurIPS workshop on Bayesian Deep Learning*, 2018.
- [13] Hankook Lee, Sung Ju Hwang, and Jinwoo Shin. Self-supervised label augmentation via input transformations. In *International Conference on Machine Learning*, pages 5714–5724. PMLR, 2020.
- [14] Juho Lee, Yoonho Lee, Jungtaek Kim, Adam Kosiorek, Seungjin Choi, and Yee Whye Teh. Set transformer: A framework for attention-based permutation-invariant neural networks. In *International Conference on Machine Learning*, pages 3744–3753. PMLR, 2019.
- [15] Juho Lee, Yoonho Lee, Jungtaek Kim, Eunho Yang, Sung Ju Hwang, and Yee Whye Teh. Bootstrapping neural processes. *arXiv preprint arXiv:2008.02956*, 2020.
- [16] Kwonjoon Lee, Subhansu Maji, Avinash Ravichandran, and Stefano Soatto. Meta-learning with differentiable convex optimization. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pages 10657–10665, 2019.
- [17] Lisha Li, Kevin Jamieson, Giulia DeSalvo, Afshin Rostamizadeh, and Ameet Talwalkar. Hyperband: A novel bandit-based approach to hyperparameter optimization. *The Journal of Machine Learning Research*, 18(1):6765–6816, 2017.
- [18] Subhansu Maji, Esa Rahtu, Juho Kannala, Matthew Blaschko, and Andrea Vedaldi. Fine-grained visual classification of aircraft. *arXiv preprint arXiv:1306.5151*, 2013.
- [19] Nikhil Mishra, Mostafa Rohaninejad, Xi Chen, and Pieter Abbeel. A simple neural attentive meta-learner. *arXiv preprint arXiv:1707.03141*, 2017.
- [20] Eunbyung Park and Junier B Oliva. Meta-curvature. *arXiv preprint arXiv:1902.03356*, 2019.

- [21] Carl Edward Rasmussen. Gaussian processes in machine learning. In *Summer School on Machine Learning*, pages 63–71. Springer, 2003.
- [22] Mengye Ren, Eleni Triantafillou, Sachin Ravi, Jake Snell, Kevin Swersky, Joshua B Tenenbaum, Hugo Larochelle, and Richard S Zemel. Meta-learning for semi-supervised few-shot classification. *arXiv preprint arXiv:1803.00676*, 2018.
- [23] Olga Russakovsky, Jia Deng, Hao Su, Jonathan Krause, Sanjeev Satheesh, Sean Ma, Zhiheng Huang, Andrej Karpathy, Aditya Khosla, Michael Bernstein, et al. Imagenet large scale visual recognition challenge. *International journal of computer vision*, 115(3):211–252, 2015.
- [24] Casper Kaae Sønderby, Tapani Raiko, Lars Maaløe, Søren Kaae Sønderby, and Ole Winther. Ladder variational autoencoders. In *Advances in neural information processing systems*, pages 3738–3746, 2016.
- [25] Yee Whye Teh, Victor Bapst, Wojciech Marian Czarnecki, John Quan, James Kirkpatrick, Raia Hadsell, Nicolas Heess, and Razvan Pascanu. Distral: Robust multitask reinforcement learning. *arXiv preprint arXiv:1707.04175*, 2017.
- [26] Oriol Vinyals, Charles Blundell, Timothy Lillicrap, Koray Kavukcuoglu, and Daan Wierstra. Matching networks for one shot learning. *arXiv preprint arXiv:1606.04080*, 2016.
- [27] Pauli Virtanen, Ralf Gommers, Travis E. Oliphant, Matt Haberland, Tyler Reddy, David Cournapeau, Evgeni Burovski, Pearu Peterson, Warren Weckesser, Jonathan Bright, Stéfan J. van der Walt, Matthew Brett, Joshua Wilson, K. Jarrod Millman, Nikolay Mayorov, Andrew R. J. Nelson, Eric Jones, Robert Kern, Eric Larson, C J Carey, İlhan Polat, Yu Feng, Eric W. Moore, Jake VanderPlas, Denis Laxalde, Josef Perktold, Robert Cimrman, Ian Henriksen, E. A. Quintero, Charles R. Harris, Anne M. Archibald, Antônio H. Ribeiro, Fabian Pedregosa, Paul van Mulbregt, and SciPy 1.0 Contributors. SciPy 1.0: Fundamental Algorithms for Scientific Computing in Python. *Nature Methods*, 17:261–272, 2020. doi: 10.1038/s41592-019-0686-2.
- [28] Catherine Wah, Steve Branson, Peter Welinder, Pietro Perona, and Serge Belongie. The caltech-ucsd birds-200-2011 dataset. 2011.
- [29] James T Wilson, Frank Hutter, and Marc Peter Deisenroth. Maximizing acquisition functions for bayesian optimization. *arXiv preprint arXiv:1805.10196*, 2018.
- [30] Chenglin Yang, Lingxi Xie, Chi Su, and Alan L Yuille. Snapshot distillation: Teacher-student optimization in one generation. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pages 2859–2868, 2019.
- [31] Huaxiu Yao, Ying Wei, Junzhou Huang, and Zhenhui Li. Hierarchically structured meta-learning. *arXiv preprint arXiv:1905.05301*, 2019.
- [32] Mingzhang Yin, George Tucker, Mingyuan Zhou, Sergey Levine, and Chelsea Finn. Meta-learning without memorization. *arXiv preprint arXiv:1912.03820*, 2019.
- [33] Jaesik Yoon, Taesup Kim, Ousmane Dia, Sungwoong Kim, Yoshua Bengio, and Sungjin Ahn. Bayesian model-agnostic meta-learning. In *Proceedings of the 32nd International Conference on Neural Information Processing Systems*, pages 7343–7353, 2018.
- [34] Linfeng Zhang, Jiebo Song, Anni Gao, Jingwei Chen, Chenglong Bao, and Kaisheng Ma. Be your own teacher: Improve the performance of convolutional neural networks via self distillation. In *Proceedings of the IEEE/CVF International Conference on Computer Vision*, pages 3713–3722, 2019.
- [35] Shengjia Zhao, Jiaming Song, and Stefano Ermon. Towards deeper understanding of variational autoencoding models. *arXiv preprint arXiv:1702.08658*, 2017.
- [36] Luisa Zintgraf, Kyriacos Shiarli, Vitaly Kurin, Katja Hofmann, and Shimon Whiteson. Fast context adaptation via meta-learning. In *International Conference on Machine Learning*, pages 7693–7702. PMLR, 2019.